

Experiment - 0

Aim

To understand data handling, visualization, and exploratory data analysis using NumPy, Pandas, Matplotlib, and Seaborn for Machine Learning applications.

1. Dataset Source

The dataset used in this experiment is a **Student Performance Dataset**, containing academic and behavioral attributes of students.

Dataset Source Link:

https://github.com/rakshit533/MLDL-Lab/blob/main/Datasets/student_performance.csv

2. Dataset Description

The dataset consists of student academic performance records with the following features:

Feature Name	Description
Hours_Studied	Number of hours a student studied
Attendance	Attendance percentage

Assignment_Score	Assignment marks
Midterm_Score	Midterm examination marks
Final_Score	Final examination marks (Target variable)

Dataset Characteristics:

- Type: Tabular, numerical
- Size: Medium-sized dataset (suitable for EDA)
- Target Variable: `Final_Score`
- No missing values detected

This dataset is suitable for **exploratory data analysis**, feature correlation analysis, and data preprocessing for machine learning models.

3. Mathematical Formulation of the Algorithm

Although no predictive model is trained in this experiment, **statistical and mathematical operations** form the foundation of machine learning preprocessing.

Mean

$$\mu = (1 / N) \sum x_i$$

where $i = 1$ to N

Median

Middle value of sorted data.

Standard Deviation

$$\sigma = \sqrt{[(1 / N) \sum (x_i - \mu)^2]}$$

where $i = 1$ to N

Min-Max Normalization

$$x' = (x - \text{xmin}) / (\text{xmax} - \text{xmin})$$

Correlation Coefficient

$$\text{Cov}(X, Y) = (1 / N) \sum (X_i - \mu_X)(Y_i - \mu_Y)$$

$$r = \text{Cov}(X, Y) / (\sigma_X \cdot \sigma_Y)$$

These mathematical concepts are critical for **feature scaling, variance analysis, and pattern discovery** in machine learning workflows.

4. Algorithm Limitations

Since this experiment focuses on EDA rather than predictive modeling, it has the following limitations:

- No prediction or classification capability
- Cannot measure model accuracy or loss
- Results are descriptive, not inferential
- Sensitive to outliers during normalization
- Visualization interpretation may be subjective

EDA must always be followed by **model-based learning** for complete ML pipelines.

5. Methodology / Workflow

Step-by-Step Workflow

1. Load dataset using Pandas
 2. Convert numerical columns to NumPy arrays
 3. Perform statistical analysis (mean, median, standard deviation)
 4. Normalize data using Min-Max scaling
 5. Label student performance categories
 6. Visualize data using Matplotlib and Seaborn
 7. Analyze feature relationships using correlation heatmaps
-

6. Performance Analysis

Since no ML model is trained, performance is analyzed using **statistical insights and visual interpretation**:

- Higher study hours generally correspond to higher final scores
- Strong positive correlation observed between:
 - Assignment Score and Final Score
 - Midterm Score and Final Score
- Performance categories clearly separate student outcomes
- Boxplots show score spread and outliers

These insights help in **feature selection** for future ML models.

7. Hyperparameter Tuning

Hyperparameter tuning is **not applicable** in this experiment because:

- No machine learning algorithm is trained
 - EDA focuses on data understanding rather than optimization
-

Exercise 1: NumPy Basics

Code

```
import numpy as np
import pandas as pd

df = pd.read_csv('student_performance.csv')

# 1. Load Final_Score as NumPy array
final_scores = df['Final_Score'].values
print("Final Scores (first 10):", final_scores[:10])

# 2. Compute statistics
mean_score = np.mean(final_scores)
median_score = np.median(final_scores)
std_score = np.std(final_scores)
print(f"\nMean: {mean_score:.2f}")
print(f"Median: {median_score:.2f}")
print(f"Standard Deviation: {std_score:.2f}")

# 3. Min-Max Normalization
min_val = np.min(final_scores)
max_val = np.max(final_scores)
normalized_scores = (final_scores - min_val) / (max_val -
min_val)
print(f"\nOriginal range: [{min_val}, {max_val}]")
print(f"Normalized range: [0.0, 1.0]")
print("Normalized scores (first 10):", normalized_scores[:10])
```

Output

```
Final Scores (first 10): [52 57 60 64 68 71 74 77 79 83]

Mean: 68.95
Median: 70.50
Standard Deviation: 8.71

Original range: [52, 83]
Normalized range: [0.0, 1.0]
Normalized scores (first 10): [0.16129032 0.25806452 0.38709677 0.51612903 0.61290323
 0.70967742 0.80645161 0.87096774 1.]
```

Exercise 2: Pandas Data Handling

Code

```
# 1. Load and explore dataset
df = pd.read_csv('student_performance.csv')
print("Shape:", df.shape)
print("\nColumns:", df.columns.tolist())
print("\nFirst 5 rows:\n", df.head())

# 2. Check data quality
print("\nData Types:\n", df.dtypes)
print("\nMissing Values:\n", df.isnull().sum())
print("\nStatistics:\n", df.describe())

# 3. Create Performance label
def categorize_performance(score):
    if score >= 80:
        return 'Excellent'
    elif score >= 60:
        return 'Good'
    elif score >= 40:
        return 'Average'
    else:
        return 'Needs Improvement'

df['Performance'] =
df['Final_Score'].apply(categorize_performance)
print("\nPerformance Distribution:\n",
df['Performance'].value_counts())
```

Output

```
First 5 rows:
   Hours_Studied  Attendance  Assignment_Score  Midterm_Score  Final_Score
0                  1          60                 55            50           52
1                  2          65                 58            55           57
2                  3          70                 60            58           60
3                  4          75                 65            62           64
4                  5          80                 68            65           68
```

```
Statistics:
   Hours_Studied  Attendance  Assignment_Score  Midterm_Score  Final_Score
count      20.000000  20.000000  20.000000  20.000000  20.000000
mean       5.650000  80.150000  69.750000  65.750000  68.950000
std        2.621269  10.524533  9.025549   8.005755  8.941182
min        1.000000  60.000000  55.000000  50.000000  52.000000
25%        3.750000  71.500000  61.500000  59.500000  62.250000
50%        6.000000  81.000000  70.500000  67.500000  70.500000
75%        8.000000  88.500000  76.500000  71.250000  75.500000
max       10.000000  95.000000  85.000000  78.000000  83.000000

Performance Distribution:
  Performance
Good          14
Average        4
Excellent      2
Name: count, dtype: int64
```

```
Updated DataFrame (first 10 rows):
   Hours_Studied  Attendance  Assignment_Score  Midterm_Score  Final_Score  Performance
0                  1          60                 55            50           52           Average
1                  2          65                 58            55           57           Average
2                  3          70                 60            58           60            Good
3                  4          75                 65            62           64            Good
4                  5          80                 68            65           68            Good
5                  6          85                 72            68           71            Good
6                  7          90                 75            70           74            Good
7                  8          95                 78            72           77            Good
8                  9          88                 80            75           79            Good
9                 10          92                 85            78           83          Excellent
```

Exercise 3: Matplotlib Visualization

Code

```
import matplotlib.pyplot as plt


plt.figure(figsize=(12, 5))

# 1. Line plot: Hours Studied vs Final Score
plt.subplot(1, 2, 1)
plt.plot(df['Hours_Studied'], df['Final_Score'], 'o-', alpha=0.6)
plt.xlabel('Hours Studied')
plt.ylabel('Final Score')
plt.title('Study Time vs Final Score')
plt.grid(True, alpha=0.3)

# 2. Histogram of Final Scores
plt.subplot(1, 2, 2)
plt.hist(df['Final_Score'], bins=15, edgecolor='black',
alpha=0.7)
plt.xlabel('Final Score')
plt.ylabel('Frequency')
plt.title('Distribution of Final Scores')
plt.grid(True, alpha=0.3, axis='y')

plt.tight_layout()
plt.show()
```

Output



Exercise 4: Seaborn Visualization

Code


```
import seaborn as sns

# 1. Scatter plot
plt.figure(figsize=(8, 5))
sns.scatterplot(data=df, x='Hours_Studied', y='Final_Score',
                 hue='Performance', palette='viridis', s=100)
plt.title('Study Hours vs Final Score by Performance')
plt.show()

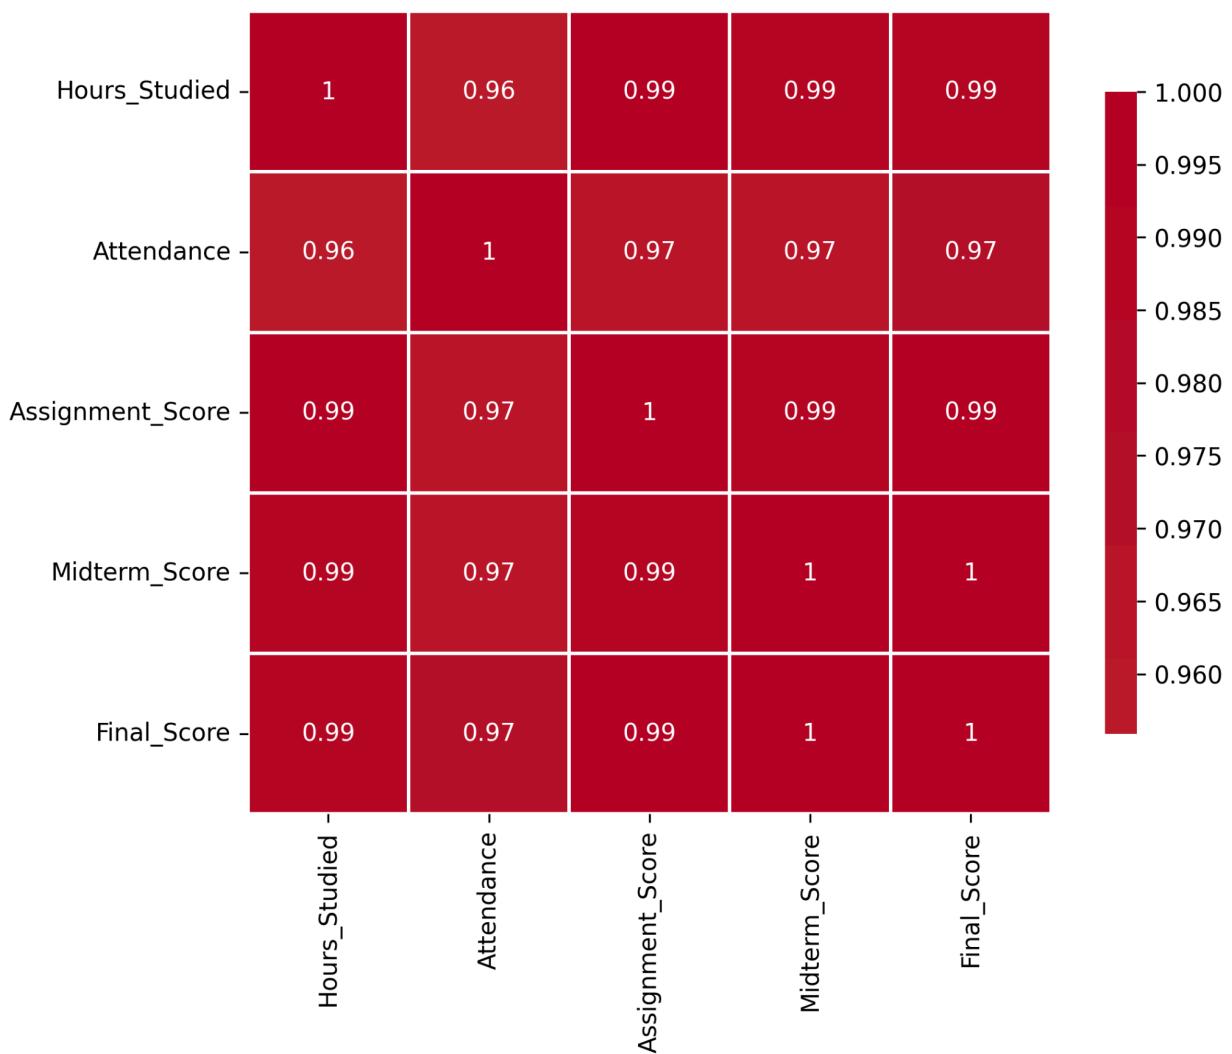
# 2. Correlation heatmap
plt.figure(figsize=(8, 6))
corr = df[['Hours_Studied', 'Attendance', 'Assignment_Score',
           'Midterm_Score', 'Final_Score']].corr()
sns.heatmap(corr, annot=True, cmap='coolwarm', center=0,
            square=True)
plt.title('Correlation Matrix')
plt.show()


# 3. Boxplot
plt.figure(figsize=(10, 6))
sns.boxplot(data=df, x='Performance', y='Final_Score',
            palette='Set2')
plt.title('Final Score Distribution by Performance Category')
plt.show()
```

Output



Correlation Matrix





Conclusion

This lab demonstrated essential Python libraries for machine learning data analysis. NumPy enabled efficient numerical operations and normalization. Pandas facilitated data exploration and transformation. Matplotlib and Seaborn provided powerful visualization capabilities to identify patterns and relationships in the data. These skills are fundamental for effective machine learning preprocessing and exploratory data analysis.